

Improving Webcam Eye-Tracking: Impact of Hardware and Head Stabilization on Data Quality

Jason Geller^{1*} , João Veríssimo²  & Julia Droulin³ 

¹Boston College, Department of Psychology and Neuroscience

²University of Lisbon, School of Arts and Humanities

³University of North Carolina - Chapel Hill, Speech and Hearing Division

Webcam-based eye-tracking offers a scalable and accessible alternative to traditional lab-based systems. While recent studies demonstrate that webcam eye-tracking can replicate canonical effects across domains such as language, memory, and decision-making, questions remain about its precision and reliability. In particular, spatial accuracy, temporal resolution, and attrition rates are often poorer than those observed with research-grade systems, raising the possibility that environmental and hardware factors introduce substantial noise. The present registered report directly tests two factors that may introduce noise into webcam data: camera quality and head stabilization. In Experiment 1, we examine the effect of external webcam quality (high vs. standard) in a single word Visual World Paradigm (VWP) task, testing whether using a better webcam can yield stronger competition effects, earlier effect onsets, and reduced attrition. In Experiment 2, we assess the impact of head stabilization (chin rest vs. no chin rest) under identical environmental conditions. Together, these studies identify the impact of hardware and movement on webcam eye-tracking data quality. Results will inform a more methodological understanding of webcam-based eye-tracking, clarifying whether its current limitations are intrinsic to the technology or can be mitigated through improved hardware and experimental control. These set of studies have implications for both online and in-lab utilization of webcam eye-tracking.

Keywords: wWebcam eye-tracking, Webcams, Visual world paradigm, Lab-based experimentation, Competition, Spoken word recognition

Words: 5806

Online experimentation in the behavioral sciences has advanced considerably since its introduction at the 1996 Society for Computers in Psychology (SCiP) conference in Chicago, IL (Reips, 2021). One methodological domain that has shown particular promise in moving online is eye tracking. Traditionally, eye-tracking studies required controlled laboratory settings equipped with specialized and costly hardware—a process that is both resource- and time-intensive. More recently, however, a growing body of research has shown that eye tracking can be successfully adapted to online environments (Bogdan et al., 2024; Bramlett & Wiener, 2024; Cruysen et al., 2023; James et al., 2025; Prystauka et al., 2024; Slim et al., 2024; Slim & Hartsuiker, 2023; Vos et al., 2022; Yang & Krajbich, 2021; Özsoy et al., 2023). By leveraging computers with webcams, researchers can now record eye movements remotely, making it possible to collect data from virtually any location at any time. This shift not only enhances scalability, but also broadens access to more diverse and representative participant samples.

Webcam-based eye tracking has become an increasingly viable and accessible method for behavioral research. Implementation typically requires only a standard computing device (e.g., laptop, desktop, tablet, or smartphone) equipped with a built-in or external webcam. Data are collected through a web browser running dedicated software capable of recording and estimating gaze position in real time.

To reliably estimate where users are looking, webcam-based eye tracking typically relies on appearance-based methods, which infer gaze direction directly from visual features of the eye region (e.g., pupil and iris appearance) (Cheng et al., 2024; Saxena et al., 2024). Recent work has extended these methods using deep learning to learn gaze-

*Send correspondence to: Jason Geller, drjasongeller@gmail.com.

Created with Quarto 1.8.27 and *preprint-typst*

appearance mappings directly from data (Kaduk et al., 2023; Saxena et al., 2024). This contrasts with research-grade eye trackers, which use model-based algorithms combining infrared illumination with geometric modeling of the pupil and corneal reflections (Cheng et al., 2024).

The most widely used library for webcam eye tracking is WebGazer.js (Papoutsaki et al., 2016; Patterson et al., 2025). WebGazer.js is an open-source JavaScript library that performs real-time gaze estimation using standard webcams. It is an appearance-based method that leverages computer vision techniques to detect the face and eyes, extract image features, and map these features onto known screen coordinates during a brief calibration procedure. Once trained, gaze locations on the screen are estimated via machine learning (Papoutsaki et al., 2016). Webgazer.js has been implemented in several experimental platforms, including Gorilla (Anwyl-Irvine et al., 2019), PsychoPy/PsychoJS (Peirce et al., 2019), jsPsych (Leeuw, 2014), PCIBex (Zehr & Schwarz, 2022), and Labvanced (Kaduk et al., 2023), making it widely accessible to researchers.

Although webcam eye-tracking is still relatively new, validation efforts are steadily accumulating and the results are encouraging. Researchers have successfully applied webcam eye-tracking with WebGazer.js to domains such as **psycholinguistics** (Bogdan et al., 2024; Buxton et al., 2024; Geller et al., 2025; Prystauka et al., 2024), judgment and decision-making (Yang & Kraljich, 2021), and memory (James et al., 2025). Collectively, this work demonstrates that webcam eye-tracking can yield interpretable and meaningful results that are comparable to those obtained with traditional lab-based systems.

However, there are several limitations associated with webgazer-based eye tracking. First, experimental effects are often smaller and noisier than those observed with research-grade eye trackers (Bogdan et al., 2024; Cruysen et al., 2023; Degen et al., 2021; Kandel & Snedeker, 2024; Slim et al., 2024; Slim & Hartsuiker, 2023). Second, relative to laboratory systems, both spatial accuracy/precision and effective temporal resolution tend to be reduced in webcam-based eye tracking.

Spatial accuracy refers to the extent to which estimated gaze positions deviate from the true point of gaze, whereas **precision** reflects the consistency of these estimates over time (Carter & Luke, 2020). In webcam-based systems, spatial accuracy often exceeds 1° of visual angle, with WebGazer.js achieving approximately 4° under laboratory conditions (Sammelmann & Weigelt, 2018). Such spatial imprecision can limit the reliable detection of brief or subtle gaze shifts, particularly in tasks that rely on small areas of interest or fine-grained gaze dynamics.

Temporal resolution in webcam-based eye tracking reflects not only the nominal sampling rate of the webcam—typically at or below 30 Hz [prystauka2024]—but also variability in sampling, execution-time delays, and the alignment of gaze samples with experimental events. Although a 30-Hz sampling rate may be sufficient to capture relatively slow eye-movement latencies, reduced temporal fidelity arising from timing variability and misalignment can impair the estimation of gaze onset, offset, and event-locked dynamics in paradigms requiring precise temporal alignment. Consistent with this, prior work has reported substantial variability in apparent effect timing across tasks and implementations, in some cases exceeding several hundred milliseconds (Sammelmann & Weigelt, 2018; Slim et al., 2024; Slim & Hartsuiker, 2023).

Together, these limitations make webcam-based eye tracking less suitable for paradigms requiring fine-grained spatial or temporal fidelity—such as tasks involving small or densely spaced areas of interest, rapid stimulus sequences, or millisecond-level event timing (James et al., 2025; Slim et al., 2024). An open question is whether these limitations primarily stem from the WebGazer.js algorithm itself or from environmental and hardware constraints—and, critically, whether improvements to experimental setup (e.g., webcam quality, head stabilization, and execution timing) can meaningfully mitigate these issues. Addressing this question is essential for determining whether webcam eye tracking should be viewed as fundamentally limited or as a method whose performance can be substantially improved through better implementation..

On the algorithmic side, recent work has demonstrated that modifications to WebGazer.js can yield substantial gains in temporal precision. Specifically, polling the sampling rate consistently and aligning timestamps to data acquisition rather than data completion markedly improves temporal resolution (James et al., 2025; Yang & Krajcich, 2021). Implementing these changes within online experiment platforms such as Gorilla and jsPsych has brought webcam-based eye-tracking closer to laboratory standards. For example, Prystauka et al. (2024) reported timing differences between lab and online of approximately 50 ms, while Geller et al. (2025) observed a 100 ms timing difference. Together, these findings indicate that at least some previously reported timing limitations reflect implementation choices rather than fundamental algorithmic constraints.

Beyond timing, several studies have compared WebGazer-based eye-tracking in controlled laboratory settings to fully remote online data collection, with mixed results. Semmelmann & Weigelt (2018) reported reduced spatial noise when WebGazer was used in the laboratory, whereas Slim et al. (2024) found largely comparable effects between lab-based and remote samples. Importantly, however, these studies were not designed to isolate specific contributors to data quality; instead, they captured broad differences between controlled and uncontrolled testing environments in which multiple factors—hardware, lighting, head movement, and participant behavior—varied simultaneously.

More targeted evidence suggests that hardware quality may nonetheless play a role in shaping webcam-based eye-tracking performance. Slim and Hartsuiker (2023) reported a positive association between sampling rate of the webcam and calibration accuracy, and Geller et al. (2025) found that participants who failed calibration more frequently were more likely to report using standard-quality built-in webcams and completing the study in suboptimal environments (e.g., natural lighting).

Beyond hardware characteristics, participant stability represents another potential contributor to data quality. In laboratory eye-tracking with research-grade systems, head movement substantially degrades spatial accuracy and precision (Hessels et al., 2014). Despite this extensive evidence, no study to date has directly examined the impact of head stabilization on webcam-based eye-tracking.

Proposed Research

To address environmental and technical sources of noise in webcam eye-tracking, we plan to bring participants into the lab to complete a Gorilla-hosted webcam task under standardized conditions. We manipulate two key factors (between subjects) across two experiments. Experiment 1 varies external webcam quality (high- vs. standard-quality external cameras). Experiment 2 varies head stabilization (i.e., chin rest vs. no chin rest). All sessions will use identical ambient lighting, fixed viewing distance, the same display/computer model, and controlled network settings. These manipulations specifically target sources of measurement noise induced by the quality of the webcam and the amount of movement by the participant.

To examine these factors, we employed a paradigm widely used in psycholinguistics—the Visual World Paradigm (VWP) (Cooper, 1974; Tanenhaus et al., 1995). The VWP has been successfully adapted for webcam-based eye tracking (Bramlett & Wiener, 2024; 2025; Geller et al., 2025; Prystauka et al., 2024). Although implementations vary across studies (Huettig et al., 2011), the version used here investigates phonemic competition, wherein item sets are typically constructed so that the display contains a target (e.g., carrot), a cohort competitor (e.g., carriage), two unrelated items that have no phonemic overlap with the target and competitor (e.g., tadpole). This configuration allows researchers to examine the dynamics of lexical competition—for instance, how phonologically similar words such as carriage (cohort effect) influence online spoken-word processing. Typically, fixations to cohort competitors persist longer or emerge earlier than fixations to unrelated distractors, reflecting transient lexical activation.

In the present study, we focus specifically on cohort competition effects in single-word spoken-word recognition using the Visual World Paradigm (VWP). Several studies (e.g., Geller et al. (2025); Slim et al. (2024)) have observed robust cohort competition effects using webcam-based eye tracking. However, effect sizes are sometimes smaller than those reported in traditional lab-based studies (e.g., Slim et al. (2024)), and effects tend to emerge later in time when measured with standard

webcams. This pattern suggests that increased measurement noise in webcam setups primarily introduces temporal delays—and in some cases attenuated effect sizes—rather than fundamentally altering or distorting the underlying dynamics of lexical competition.

The current research aims to inform best practices for webcam-based eye tracking, with particular attention to hardware quality and physical setup considerations (e.g., head movement). Reducing noise by manipulating hardware and head movement is predicted to make the measured gaze signal more stable and less variable across time and trials. In turn, this can make existing effects easier to detect, potentially manifesting as (a) larger and clearer competition effects, (b) earlier and more reliable detection of effect onsets in time-course analyses, and (c) lower calibration failure and attrition rates compared to standard webcams.

While these guidelines will benefit researchers conducting webcam studies in uncontrolled, online settings, they are also valuable for laboratory-based research in which webcams may serve as lower-cost alternatives to infrared eye-tracking systems. By systematically testing the role of hardware and head stabilization, this work clarifies the conditions under which webcam eye tracking can approximate lab-quality data and where its limitations remain.

Experiment 1: High-Quality Webcam vs. Standard-Quality Webcam

Both [Slim & Hartsuiker \(2023\)](#) and [Geller et al. \(2025\)](#) observed a clear relationship between webcam quality and calibration accuracy in webcam-based eye-tracking. Building on these findings, Experiment 1 tests how webcam quality influences competition effects in a single-word VWP. Specifically, we ask whether a higher-quality webcam yields (a) a greater proportion of looks to relevant interest areas (i.e., greater looks to cohorts vs. unrelated items) (b) an earlier emergence of these effects over time, and (c) lower data attrition rates relative to a lower-quality webcam.

To address this, participants will complete the same VWP task using one of two webcam types: a high-quality external webcam (e.g., Logitech Brio) and a standard external webcam designed to emulate a typical built-in laptop camera (e.g., Logitech C270). The high-quality webcam offers higher resolution, a higher sampling rate (60 Hz), and greater frame-rate stability, and more consistent illumination handling—factors expected to enhance gaze precision and tracking reliability. In contrast, more standard webcams, while representative of most participants' home setups, typically provide lower frame rates and exhibit greater variability under different lighting conditions. Comparing these two setups enables a direct assessment of how hardware quality constrains the strength, timing, and reliability of linguistic competition effects in webcam-based eye-tracking.

Hypotheses

We hypothesize several effects related to competition, onset, and attrition:

Competition Effects

(H1a) Participants will show an overall greater proportion of looks to competitors. (H1b) Webcam quality (high vs. standard) will influence the overall proportion of looks, with higher-quality webcams detecting a greater proportion of looks.

To quantify the effect of webcam quality on the proportion of looks within the time window of interest, we will use Cohen's h —a standardized measure of effect size appropriate for comparisons between two proportions.

Timing/Onset Effects

(H1b) Each webcam condition will show a change in the proportion of looks across time. More specifically we hypothesize the proportion change will be non-linear across time and that there will be a difference between the two conditions across time. (H2b) For standard-quality webcams vs. high-quality webcam, onsets will be detected later due to increased noise.

Attrition

(H1c) Attrition rates will be lower in the high-quality webcam condition than in the standard-quality webcam condition.

Method

All stimuli (audio and images), code, and data (raw and summary) will be placed on OSF at this link: <https://osf.io/cf6xr/overview>. The entire experiment will be stored on Gorilla's open materials with a link to preview the tasks. In addition, the code and manuscript will be fully reproducible using Quarto and the package manager nix (Dolstra & contributors, 2023) in combination with the R package {rix} (Rodrigues & Baumann, 2025). Together, nix and {rix} enable reproducible computational environments at both the system and package levels. This manuscript and all of the necessary files to reproduce it will be stored on GitHub. **Participants will provide informed consent prior to participation. At the time of data collection, all study procedures have been approved by the relevant ethics committee.**

Sampling Goal

We conducted an *a priori* power analysis using the {mixedPower} package [kumle2021] in R. Data from 21 participants, collected online using the Gorilla experimental platform during development of the {webgazeR} package [gellerLanguageBordersStepbystep2025f] and employing the same stimuli and visual-world word paradigm (VWWP) design, were used to seed the simulations. These pilot data did not manipulate webcam quality. In these data, we observed a cohort effect of approximately 3%.

Using this value as the effect-size seed, we collapsed the data across time bins to compute binomial fixation counts per trial and fit a binomial generalized linear mixed-effects model (GLMM) to obtain fixed-effect estimates. Because the planned experiment included a manipulation of webcam quality, we conducted simulations that incorporated this factor. Specifically, we augmented the dataset with a hypothetical between-subjects factor for webcam quality, with participants evenly assigned to high-quality and standard-quality webcam groups. In the high-quality webcam group, we modeled both a higher overall fixation rate and a larger cohort effect; in the standard-quality webcam group, the cohort effect was reduced by half relative to the high-quality group.

Simulated datasets were generated under this model, and the planned GLMM—including a condition $\times$ webcam-quality interaction—was refit to each simulated dataset (5,000 simulations). Statistical power was estimated as the proportion of simulations in which the interaction term exceeded $|z| = 1.96$. The full analysis script used to conduct this power analysis is available on OSF (<https://osf.io/4trmn/files/a46g8>).

Results indicated that a sample size of 35 participants per group ($N = 70$ total) would provide approximately 90% power to detect the hypothesized reduction in the cohort effect and overall fixation rate under standard-quality webcam conditions. Because this analysis focused on overall fixation proportions and we additionally plan to examine effects across time, we will recruit 50 participants per group ($N = 100$ total). Data collection will continue until 100 usable participants are obtained (50 per webcam-quality group). That is, participants will be replaced if they fail calibration at any point. For analyses of attrition (i.e., failed calibration attempts) (see below), all participants who enter the study will be included, thus we might have more than 100 participants total.

Materials

VWP

Picture Stimuli. Stimuli were adapted from @colby2023. Each stimulus set comprised four images. For the webcam study, we used 60 stimulus sets (30 monosyllabic, 30 bisyllabic). All trials were of the TCUU type (target-cohort-unrelated-unrelated), in which a target, its onset cohort competitor, and two unrelated images were displayed (e.g., MOUTH, mouse, house, chain). This design yielded 60 trials in total, with each stimulus set contributing one trial. A custom MATLAB script (R2024a; <https://osf.io/x3frv>) generated a unique randomized trial list for each participant, pseudo-randomizing display positions such that the target, cohort, and unrelated images were approximately equally likely to appear in each quadrant across participants.

All 120 images were drawn from a commercial clipart database, selected by a small focus group of students, and edited using a standard lab protocol to ensure a cohesive visual style [mcmurray2010]. All images were scaled to 300 × 300 pixels.

Auditory Stimuli. Auditory stimuli were recorded by a female monolingual speaker of English in a sound-attenuated room sampled at 44.1 kHz. Auditory tokens were edited to reduce noise and remove clicks. They were then amplitude normalized to 70 dB SPL. All .wav files were converted to .mp3 for online data collection.

Webcams

To manipulate recording quality, two webcams will be used. In the high-quality condition, we will use a Logitech Brio webcam, which records in 4K resolution (up to 4096 × 2160 px) with a 90° field of view and samples at 60 Hz. This setup provides high-fidelity video with greater spatial and temporal precision. In the standard-quality condition, we will use a Logitech C270 HD webcam, which records in 720p resolution and samples at 30 Hz, producing video quality comparable to that of a typical built-in laptop webcam and therefore simulating lower-quality online recordings (Jarvis et al., 2025).

Both webcams will be mounted in a fixed position above the monitor to maintain consistent framing across participants. Lighting will be standardized to ensure uniform image quality across all sessions.

Experimental Setup and Procedure

All tasks were completed in a single session lasting approximately 30 minutes. The experiment was programmed and administered in Gorilla (Anwyl-Irvine et al., 2020) **and run in the Google Chrome browser**. Participants were tested in a dedicated room in the Human Neuroscience Lab at Boston College with **two computers** and were seated approximately 65 cm from a 23-inch Dell U2312HM monitor (1920 × 1080 px). Each testing computer was connected to the same network via a wired ethernet connection (**internet speed was tested before each session to ensure a high-speed connection**). Auditory information was presented over Sony ZX110 headphones to ensure consistent audio delivery and to minimize background noise.

Experimental tasks were fixed and presented in the following order: informed consent, the single-word visual world paradigm (VWP) task, and a demographic questionnaire. The complete experiment is available for viewing on the Gorilla platform.

Webcam Single Word VWP

Webcam Calibration. To ensure that participants understood how webcam-based eye tracking works, they first viewed an instructional video demonstrating the calibration procedure (available on OSF here: [link]). Calibration was completed twice—once at the beginning of the experiment and again after 30 trials. Each calibration session allowed up to three attempts.

During calibration, participants were asked to position themselves so that a face mesh could be successfully fit to their face. Once done, participants completed a passive calibration task in which nine red targets were presented sequentially on the screen. Participants were instructed to look directly at each target while it was visible. We used default system parameters for the calibration procedure, including the required fixation duration per target (200 ms), the number of gaze samples used per prediction (10 prediction points), and the transition time before data collection began (1000 ms). Immediately following calibration, validation was performed using five green targets presented one at a time. Participants were again instructed to look directly at each target. During validation, the eye-tracking system evaluated calibration quality by comparing the predicted gaze location to the known target location. Calibration was considered unsuccessful if more than two validation points deviated from their intended target location (i.e., predicted gaze was misaligned with another calibration point). Participants who failed any calibration check were branched to the end of the experiment and asked to complete a brief demographic questionnaire.

VWP. After calibration, participants will then complete four practice trials to familiarize themselves with the task. Each trial begins with a 500 ms central fixation cross, followed by a

preview display of four images located in the screen's corners. After 1500 ms, a start button appears at the center; participants click it to confirm fixation before hearing the spoken word. The images remain visible throughout the trial, and participants indicate their response by clicking the image corresponding to the spoken target. A response deadline of 5 seconds will be used. **Eye movements will be recorded continuously during the final image display.** Please see Figure 1 for a schematic of the trial. After the main task, participants complete a brief demographic questionnaire and are thanked for their participation.




Figure 1. Trial Schematic for webcam VWP.

Demographic Questionnaire

After the main task, we will have participants complete a demographic questionnaire. The questions cover basic demographic information, including age, gender, spoken dialect, ethnicity, and race.

Data Preprocessing and Exclusions

We will follow the guidelines outlined in Geller et al. (2025) and exclude participants with overall task accuracy below 80%, those who report English as not their first language, and those with non-normal or uncorrected vision. At the trial level, only correct-response trials (accuracy = 1) will be retained. Reaction times (RTs) lower or greater than 2.5 SD by-participant will be removed.

For eye-tracking preprocessing we will use the {webgazer} package in R (Geller et al., 2025) that contains helper functions to preprocess webcam eye-tracking data. All webcam eye-tracking files and behavioral data will be merged. Data quality will be screened via sampling-rate checks with very low-frequency recordings (e.g., < 30 Hz)¹ by-participant and by-trial excluded (Bramlett & Wiener, 2025; Vos et al., 2022). We will quantify out-of-bounds (OOB) samples—gaze points outside the normalized screen (1,1)—and remove participants and trials with excessive OOB data (> 30%). OOB samples will be discarded prior to analysis. **In addition, Gorilla provides two eye-tracking quality metrics derived from the underlying face-tracking model: convergence and confidence. Convergence (range: 0–1) reflects the model’s certainty that a face has been successfully detected, with higher values indicating poorer convergence on a face. Confidence (range: 0–1) reflects the support vector machine (SVM) classifier’s confidence in the detected face. Trials will be excluded if convergence exceeds 0.5 (indicating unreliable face detection) or if confidence falls below 0.5 (indicating low classifier certainty). To increase signal-to-noise, participants with fewer than 40 usable trials after these exclusions will also be removed.**

¹Given inconsistent sampling in webgazer if too many participants and trials are thrown out with less than 30 hz we will lessen this criteria to 15 hz.

Areas of Interest (AOIs) will be defined in normalized coordinates as the four screen quadrants, and gaze samples will be assigned to AOIs. Trial time will be aligned to the actual stimulus onset by taking the audio onset metric provided by Gorilla. We then subtract 100 ms due to silence prefixed to the audio recording.

Analysis Plan

GAMMs

For analysis, gaze data will be segmented into 50-ms time bins². A look is defined as any time point in which the WebGazer.js algorithm returns a valid gaze estimate. Within each time bin, gaze samples falling on the target item were excluded. The remaining samples were classified as directed toward either the cohort competitor or an unrelated competitor, with looks to either unrelated item collapsed into a single unrelated category. For each participant $\times$ trial $\times$ time bin, we counted the number of gaze samples falling on the cohort and on unrelated items. Each bin was then coded as cohort-dominant (1) if the cohort received more samples than the unrelated category, and as 0 otherwise. These binary values were subsequently summed across trials, yielding, for each participant $\times$ time bin, the number of trials in which gaze favored the cohort vs. unrelated looks. This cohort-dominance count served as the dependent measure for statistical modeling and summary analyses.

This simplified analysis approach offers several advantages. First, aggregating gaze data into binomial outcomes substantially reduces data dimensionality and mitigates issues related to temporal autocorrelation that are common in high-frequency eye-tracking data (Veríssimo & Lago, 2025). Second, within a generalized additive mixed-model (GAMM) framework, using a binomial response allows temporal autocorrelation to be modeled directly, which we implement in the analyses described below.

To analyze overall competition effects and onset latency, we will use generalized additive mixed models [GAMMs; (Wood, 2017)]. GAMMs extends the generalized linear modeling framework by modeling effects that are expected to vary non-linearly over time—a common feature in the VWP (Brown-Schmidt et al., 2025; Ito & Knoeferle, 2022; Mitterer, 2025; Veríssimo & Lago, 2025). These models capture non-linear effects by fitting smoothing splines to the data using data-driven, machine-learning-based methods, with the amount of non-linearity captured by how “wiggly” the time course is (one can think of this as the number of bowpoints of a curve). A benefit of this approach is that it reduces the risk of over-fitting and eliminates the need to use polynomial terms, as required in traditional growth curve models (Mirman, 2014). Importantly, GAMMs also allow researchers to account for autocorrelation in time-series data, which is especially critical in gaze analyses where successive samples are not independent. By modeling the autocorrelation structure, GAMMs provide more accurate estimates of temporal effects and prevent inflation of Type I error rates (Rij et al., 2019). In addition to this, fitting GAMMs allow us to estimate the onset of the competition effect in each condition (Veríssimo & Lago, 2025).

Competition Effects (H1a). Gaze samples will be analyzed with a binomial (logistic) GAMM using the `bam()` function from the `{mgcv}` package (Wood, 2017). For visualization, we will employ functions from the `{tidygam}` package (Coretta, 2024), the `{onsets}` package (Veríssimo & Lago, 2025) to examine onset latencies, and the `{itsadug}` (van Rij et al., 2022) package for AR functions and to test differences in smooth splines across the time (`get_differences()`). The dependent variable will consist of gaze counts to cohorts compared to unrelated items, for each participant and in each 100ms time bin. All analyses were conducted on a window ranging from stimulus onset (100 ms) to 1200 ms.

The right model in the context of GAMMs is a difficult one. Results have been shown to vary depending on whether one uses an ordered factor scheme or unordered factor scheme with time (Oltrogge et al., 2025). Because of this we plan to fit the model multiple ways to test for robustness. In one model (Listing 2) we will fit a model that includes a parametric term for webcam type that is an unordered factor (treatment-coded such that high-quality = 1 and standard-

²We might increase this to 100 ms if too many empty bins.

quality = 0). To examine whether webcam type moderates the cohort effect over time, we will include smooth terms for time-by-condition interactions with camera. To account for individual differences, we will include participant-specific random smooths for time, and participant-specific random smooths for time within each level of the camera factor.

This specification allows the model to capture three components:

1. The overall (parametric) effect of webcam type on the proportion of looks to cohorts-over-unrelated items. Here, the intercept reflects the expected proportion of looks in the standard-quality condition, and the webcam coefficient reflects the difference in looks between high- and standard-quality webcams.
2. Condition-specific, time-varying trajectories in cohort competition (via the smooth terms), which support inferences about the timing (onsets) and dynamics of cohort effects across conditions. With an unordered factor, the model estimates time smooths for each condition separately with the p -value telling us if each smooth differs from zero.
3. Deviations from these group-level trajectories at the individual participant level (via random smooths)

In a second instantiation of the model (see Listing 3), we will include a parametric term for webcam type treated as an ordered factor. To assess whether webcam type moderates the cohort-over-unrelated effect over time, we will incorporate smooth terms for time as well as time-by-condition interactions, as required when fitting ordered factors.

This specification allows the model to capture the same effects above with one key difference: By including an ordered factor we must include separate smooths for time (the reference) and time $\times$ webcam type. This model estimates a baseline smooth for the standard-quality webcam and a difference smooth for the high-quality webcam. Because the webcam factor is ordered, the difference smooth directly represents how the trajectory for high-quality deviates from the standard-quality trajectory. This structure enables inferences about the timing and dynamics of cohort effects across webcam conditions. We will also look at how changing the order affects inferences.

For both models, we use default arguments for the smooth functions. We use $k = 9$ (number of knots) to capture the “wigglyness” of the smooths. Smooths were fit with a thin plate regression spline. This default seems to work generally well (Oltrogge et al., 2025); however deviations from this default will be reported and robustness will be tested.

Although “maximal” random-effects structures are often recommended in linear mixed-effects models (Barr et al., 2013), such specifications can be computationally prohibitive in GAMMs. The present specification follows the recommendations of Verissimo & Lago (2025). Statistical significance will be assessed at $\alpha = .05$. When fitting difference curves with ordered factors with simultaneous CIs, we will correct for multiple comparisons using a Bonferroni correction (Krause et al., 2024).

To account for autocorrelation in the residuals, we will first fit the model without an autoregressive term in order to estimate the autocorrelation parameter (ρ). We will then re-fit the model including a first-order autoregressive process (AR(1)) to properly model temporal dependencies. Although using larger time bins can reduce autocorrelation, it does not eliminate it entirely, so explicitly modeling residual autocorrelation ensures valid statistical inference.

Onset Effects (H1b). Once the models are fit, we will extract predicted gaze curves for each condition using the {onsets} package (Verissimo & Lago, 2025). The {onsets} procedure first simulates gaze curves across time from the fitted GAMMs ($N = 10,000$). For each simulated curve, the onset of the condition effect is identified by comparing the predicted log-odds at each timepoint to a predefined criterion. Within the package, onset is defined as the earliest time at which the predicted log-odds is significantly greater than the model-predicted log-odds at the first timepoint of the analysis window (here word onset). Repeating this procedure across 10,000 simulations yields a distribution of onset estimates, from which a 95% highest density interval (HDI) can be obtained. To derive between-condition comparisons, onset times from paired simulations are subtracted, producing a corresponding distribution with median onset difference and its associated 95% HDI (see Listing 4 for the analysis code).

Listing 1. R packages to load

```
# load packages
library(mgcv) # perform gam
library(tidygam) # visualization of gams
library(itsadug) # get_differences
library(onsets) # onset differences
```

Listing 2. GAMM model set-up unordered camera factor

```
# set contrasts
options(contrasts = rep("contr.treatment", 2)) # treatment code variables

# combine levels of both factors into one factor
dat$camera <- as.factor(dat$camera)

# get start event for AR
start_event = start_event(., column = "Time", # for ar
                           event = c("subj", "Trial"),
                           label = "start.event",
                           label.event = NULL,
                           order = FALSE)$start.event)

# quick rho estimate (fit once without AR to get residual ACF ~ lag1)

m0 <-
bam(cbind(fix_cohort, fix_unrelated) ~ 1 +
    camera + s(time, by = camera, k = 10) +
    s(participant, by=camera, bs = "re") +
    s(time, participant, by=camera, bs = "re"),
    family = binomial(), method = "fREML",
    discrete = TRUE, data = dat, na.action = na.omit)

rho <- acf(residuals(m0, type = "pearson"), plot = FALSE)$acf[2]

# final model with AR(1) to handle within-series autocorrelation

m1 <-
bam(cbind(fix_cohort, fix_unrelated) ~ 1 +
    camera + s(time, by = camera, k = 10) +
    s(participant, by=camera, bs = "re") +
    s(time, participant, by=camera, bs = "re") +
    rho=rho, AR.start = start_event)
```


Listing 5. Getting onset differences using the {onsets} package for ordered model

```
# Obtain onsets in each condition (and their differences)

onsets_comp <- get_onsets(model = m1,           # Fitted GAMM
                          time_var = "time",   # Name of time variable
                          by_var = "camera_ord", # Name of condition/group variable
                          compare = T,         # Obtain differences

  between onsets
                          n_samples = 10000,    # Large number of
  samples (less variable results)
                          seed = 1)            # Random seed for
  reproducibility
```

Attrition (Calibration Failure) (H1c)

To examine whether webcam type affects calibration failure and thus removal from the study, we will fit a logistic regression model using the `glm()` function (see Listing 6). Calibration outcome will be coded as a binary variable, where 0 indicates that a participant failed to calibrate at either of the calibration phases, and 1 indicates that the participant successfully passed both calibration phases. This model will allow us to estimate whether the webcam condition reliably predicts the probability of successful calibration.

Listing 6. Logistic model to examine the effect of webcams on calibration

```
# fit glm model
glm(calibration ~ camera, family = binomial(link = "logit"))
```

Experiment 2: Head stabilization (chin rest) vs. no head stabilization (no chin rest)

In Experiment 2, we use the same standard-quality external webcam as in Experiment 1 but manipulate head stability by comparing a chin-rest condition with a no-chin-rest condition. Prior work using laboratory-based eye-tracking systems has shown that head movement can substantially degrade the accuracy and reliability of gaze estimates. For example, @hessels2014 demonstrated that even moderate deviations in head position can produce systematic calibration drift, increased data loss, and slower recovery following tracking loss. Importantly, these effects are not restricted to extreme movements and can arise during typical participant behavior when head position is unconstrained.

Although some online platforms [e.g., LabVanced; Finger et al. (2017)] offer proprietary head-tracking or motion-monitoring solutions, other commonly used platforms—such as PsychoPy and Gorilla—require custom code to implement comparable functionality when using WebGazer.js. At present, it remains unclear how warning-based motion-control approaches interact with WebGazer’s estimates of event detection, onset latency, and gaze competition.

In laboratory-based eye-tracking, chin rests are routinely used to stabilize head position so that gaze estimates primarily reflect eye movements rather than head motion. While chin rests are uncommon in fully remote testing, introducing a chin-rest manipulation in the present study provides a controlled and conservative test of the causal role of head stability in webcam-based eye-tracking. Specifically, the chin-rest condition serves as a theoretically motivated proxy for instruction-based or software-based head-stabilization approaches and allows us to estimate an upper bound on the benefits that head-movement control may provide for calibration quality, competition effects, event-detection timing, and participant attrition.

Hypotheses

The hypotheses for Experiment 2 are identical to those of Experiment 1 and are as follows:

Competition Effects (H2a)

Participants will show an overall greater proportion of looks to competitors. (H1b) Head stabilization (chin rest vs. no chin rest) will influence the overall proportion of looks, with chin rest condition detecting a greater proportion of looks.

Timing/Onset Effects (H2b)

Each condition will show a change in the proportion of looks across time. More specifically we hypothesize the proportion change will be non-linear across time and that there will be a difference between the two conditions across time. (H2b) For the no chin-rest condition compared to the chin-rest condition, onsets will be detected later due to increased noise.

Attrition (Calibration failure) (H2c)

Attrition due to calibration failures will be lower in the high-quality webcam condition than in the standard-quality webcam condition.

Sampling Goal, Materials , Procedure, and Analysis Plan

The sampling goal, materials, procedure, and analysis plan are the same as Experiment 1. The main difference is whether participants use a chin rest or not.

Declarations

Funding

This study was funded by the main author.

Conflicts of Interest

The authors declare no conflicts of interest.

Ethics Approval

This study was approved by the relevant ethics committee.

Consent for Publication

I consent for my paper to be published in BRM.

Consent to Participate

All participants will provide informed consent prior to participation.

Availability of Data, Code, and Materials

All data, code, and materials will be stored on OSF (<https://osf.io/cf6xr/overview>) and GitHub (<https://github.com/jgeller112/Webcam2Lab-VWP>).

Authors' Contributions

- JG - Writing—original draft
- JV - Writing—review & editing
- JD - Writing—review & editing

References

- Anwyl-Irvine, A. L., Massonnié, J., Flitton, A., Kirkham, N., & Evershed, J. K. (2019). Gorilla in our midst: An online behavioral experiment builder. *Behavior Research Methods*, 52(1), 388–407. <https://doi.org/10.3758/s13428-019-01237-x>
- Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013). Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, 68(3), 255–278. <https://doi.org/10.1016/j.jml.2012.11.001>
- Bogdan, P. C., Dolcos, S., Buetti, S., Lleras, A., & Dolcos, F. (2024). Investigating the suitability of online eye tracking for psychological research: Evidence from comparisons with in-person data using emotion–attention interaction tasks. *Behavior Research Methods*, 56(3), 2213–2226. <https://doi.org/10.3758/s13428-023-02143-z>
- Boxtel, W. S. van, Linge, M., Manning, R., Haven, L. N., & Lee, J. (2024). Online Eye Tracking for Aphasia: A Feasibility Study Comparing Web and Lab Tracking and Implications for Clinical Use. *Brain and Behavior*, 14(11). <https://doi.org/10.1002/brb3.70112>
- Bramlett, A. A., & Wiener, S. (2024). The art of wrangling. *Linguistic Approaches to Bilingualism*. <https://doi.org/https://doi.org/10.1075/lab.23071.bra>
- Bramlett, A. A., & Wiener, S. (2025). The Art of Wrangling: Working with Web-Based Visual World Paradigm Eye-Tracking Data in Language Research. *Linguistic Approaches to Bilingualism*, 15(4), 538–570. <https://doi.org/10.1075/lab.23071.bra>
- Brown-Schmidt, S., Cho, S.-J., Fenn, K. M., & Trude, A. M. (2025). Modeling spatio-temporal patterns in intensive binary time series eye-tracking data using Generalized Additive Mixed Models. *Brain Research*, 1854, 149511. <https://doi.org/10.1016/j.brainres.2025.149511>
- Carter, B. T., & Luke, S. G. (2020). Best practices in eye tracking research. *International Journal of Psychophysiology*, 155, 49–62. <https://doi.org/10.1016/j.ijpsycho.2020.05.010>
- Cheng, Y., Wang, H., Bao, Y., & Lu, F. (2024, April). *Appearance-Based Gaze Estimation With Deep Learning: A Review and Benchmark* (Issue arXiv:2104.12668). <https://doi.org/10.48550/arXiv.2104.12668>
- Cooper, R. M. (1974). The control of eye fixation by the meaning of spoken language: A new methodology for the real-time investigation of speech perception, memory, and language processing. *Cognitive Psychology*, 6(1), 84–107. [https://doi.org/10.1016/0010-0285\(74\)90005-X](https://doi.org/10.1016/0010-0285(74)90005-X)
- Coretta, S. (2024). *tidygam: Tidy Prediction and Plotting of Generalised Additive Models*. <https://github.com/stefanocoretta/tidygam>
- Cruyssen, I. Van der, Ben-Shakhar, G., Pertzov, Y., Guy, N., Cabooter, Q., Gunschera, L. J., & Verschuere, B. (2023). The validation of online webcam-based eye-tracking: The replication of the cascade effect, the novelty preference, and the visual world paradigm. *Behavior Research Methods*. <https://doi.org/10.3758/s13428-023-02221-2>
- Degen, J., Kursat, L., & Leigh, D. D. (2021). Seeing is believing: Testing an explicit linking assumption for visual world eye-tracking in psycholinguistics. *Proceedings of the Annual Meeting of the Cognitive Science Society*, 43(43).
- Dolstra, E., & contributors, T. N. (2023,). Nix. <https://nixos.org/>
- Finger, H., Goeke, C., Diekamp, D., Standvoß, K., & König, P. (2017,). LabVanced: A Unified JavaScript Framework for Online Studies. *International Conference on Computational Social Science (Ic2s2)*.
- Geller, J., Prystauka, Y., Colby, S. E., & Drouin, J. R. (2025). Language without Borders: A Step-by-Step Guide to Analyzing Webcam Eye-Tracking Data for L2 Research. *Research Methods in Applied Linguistics*, 4(3), 100226. <https://doi.org/10.1016/j.rmal.2025.100226>
- Huetting, F., Rommers, J., & Meyer, A. S. (2011). Using the visual world paradigm to study language processing: a review and critical evaluation. *Acta Psychologica*, 137(2), 151–171. <https://doi.org/10.1016/j.actpsy.2010.11.003>
- Ito, A., & Knoeferle, P. (2022). Analysing data from the psycholinguistic visual-world paradigm: Comparison of different analysis methods. *Behavior Research Methods*, 55(7), 3461–3493. <https://doi.org/10.3758/s13428-022-01969-3>

- James, A. N., Ryskin, R., Hartshorne, J. K., Backs, H., Bala, N., Barcenás-Meade, L., Bhattarai, S., Charles, T., Copoulos, G., Coss, C., Eisert, A., Furuhashi, E., Ginell, K., Guttman-McCabe, A., Harrison, E., Hoban, L., Hwang, W. A., Iannetta, C., Koenig, K. M., ... de Leeuw, J. R. (2025). What Paradigms Can Webcam Eye-Tracking Be Used For? Attempted Replications of Five Cognitive Science Experiments. *Collabra: Psychology*, 11(1), 140755. <https://doi.org/10.1525/collabra.140755>
- Jarvis, M., Vasarhelyi, A., Anderson, J., Mulley, C., Lipp, O. V., & Ney, L. J. (2025). js-mEye: An extension and plugin for the measurement of pupil size in the online platform jsPsych. *Behavior Research Methods*, 58(1). <https://doi.org/10.3758/s13428-025-02901-1>
- Kaduk, T., Goeke, C., Finger, H., & König, P. (2023). Webcam eye tracking close to laboratory standards: Comparing a new webcam-based system and the EyeLink 1000. *Behavior Research Methods*, 56(5), 5002–5022. <https://doi.org/10.3758/s13428-023-02237-8>
- Kandel, M., & Snedeker, J. (2024). Assessing two methods of webcam-based eye-tracking for child language research. *Journal of Child Language*, 52(3), 675–708. <https://doi.org/10.1017/s0305000924000175>
- Krause, J., Rij, J. van, & Borst, J. P. (2024). Word Type and Frequency Effects on Lexical Decisions Are Process-dependent and Start Early. *Journal of Cognitive Neuroscience*, 36(10), 2227–2250. https://doi.org/10.1162/jocn_a_02214
- Leeuw, J. R. de. (2014). jsPsych: A JavaScript library for creating behavioral experiments in a Web browser. *Behavior Research Methods*, 47(1), 1–12. <https://doi.org/10.3758/s13428-014-0458-y>
- Mirman, D. (2014). *Growth Curve Analysis and Visualization Using R*. Chapman, Hall/CRC.
- Mitterer, H. (2025). A web-based mouse-tracking task for early perceptual language processing. *Behavior Research Methods*, 57(11). <https://doi.org/10.3758/s13428-025-02827-8>
- Oltrogge, E., Verissimo, J., Patil, U., & Lago, S. (2025). Memory retrieval and prediction interact in sentence comprehension: An experimental evaluation of a cue-based retrieval model. *Journal of Memory and Language*, 144, 104651. <https://doi.org/10.1016/j.jml.2025.104651>
- Papoutsaki, A., Sangkloy, P., Laskey, J., Daskalova, N., Huang, J., & Hays, J. (2016, July). WebGazer: Scalable Webcam Eye Tracking Using User Interactions. *International Joint Conference on Artificial Intelligence*.
- Patterson, A. S., Nicklin, C., & Vitta, J. P. (2025). Methodological recommendations for webcam-based eye tracking: A scoping review. *Research Methods in Applied Linguistics*, 4(3), 100244. <https://doi.org/10.1016/j.rmal.2025.100244>
- Peirce, J., Gray, J. R., Simpson, S., MacAskill, M., Höchenberger, R., Sogo, H., Kastman, E., & Lindeløv, J. K. (2019). PsychoPy2: Experiments in behavior made easy. *Behavior Research Methods*, 51(1), 195–203. <https://doi.org/10.3758/s13428-018-01193-y>
- Prystauka, Y., Altmann, G. T. M., & Rothman, J. (2024). Online eye tracking and real-time sentence processing: On opportunities and efficacy for capturing psycholinguistic effects of different magnitudes and diversity. *Behavior Research Methods*, 56(4), 3504–3522. <https://doi.org/10.3758/s13428-023-02176-4>
- Reips, U.-D. (2021). Web-based research in psychology: A review. *Zeitschrift Für Psychologie*, 229(4), 198–213. <https://doi.org/10.1027/2151-2604/a000475>
- Rij, J. van, Hendriks, P., Rijn, H. van, Baayen, R. H., & Wood, S. N. (2019). Analyzing the Time Course of Pupillometric Data. *Trends in Hearing*, 23. <https://doi.org/10.1177/2331216519832483>
- Rodrigues, B., & Baumann, P. (2025). *rix: Reproducible Data Science Environments with 'Nix'*. <https://docs.ropensci.org/rix/>
- Saxena, S., Fink, L. K., & Lange, E. B. (2024). Deep Learning Models for Webcam Eye Tracking in Online Experiments. *Behavior Research Methods*, 56(4), 3487–3503. <https://doi.org/10.3758/s13428-023-02190-6>
- Semmelmann, K., & Weigelt, S. (2018). Online webcam-based eye tracking in cognitive science: A first look. *Behavior Research Methods*, 50(2), 451–465. <https://doi.org/10.3758/s13428-017-0913-7>
- Slim, M. S., & Hartsuiker, R. J. (2023). Moving visual world experiments online? A web-based replication of Dijkgraaf, Hartsuiker, and Duyck (2017) using PCIBex and WebGazer.js. *Behavior Research Methods*, 55(7), 3786–3804. <https://doi.org/10.3758/s13428-022-01989-z>

- Slim, M. S., Kandel, M., Yacovone, A., & Snedeker, J. (2024). Webcams as Windows to the Mind? A Direct Comparison Between In-Lab and Web-Based Eye-Tracking Methods. *Open Mind*, 8, 1369–1424. https://doi.org/10.1162/opmi_a_00171
- Tanenhaus, M. K., Spivey-Knowlton, M. J., Eberhard, K. M., & Sedivy, J. C. (1995). Integration of visual and linguistic information in spoken language comprehension. *Science (New York, N.Y.)*, 268(5217), 1632–1634. <http://www.ncbi.nlm.nih.gov/pubmed/7777863>
- van Rij, J., Wieling, M., Baayen, R. H., & van Rijn, H. (2022,). *itsadug: Interpreting Time Series and Autocorrelated Data Using GAMMs*.
- Veríssimo, J., & Lago, S. (2025, April 29). *A novel method for detecting the onset of experimental effects in visual world eye-tracking*. PsyArXiv. https://doi.org/10.31234/osf.io/yk4xb_v3
- Vos, M., Minor, S., & Ramchand, G. C. (2022). Comparing infrared and webcam eye tracking in the Visual World Paradigm. *Glossa Psycholinguistics*, 1(1). <https://doi.org/10.5070/G6011131>
- Wood, S. N. (2017). *Generalized Additive Models: An Introduction with R* (2nd ed.). Chapman, Hall/CRC.
- Yang, X., & Krajbich, I. (2021). Webcam-Based Online Eye-Tracking for Behavioral Research. *Judgment and Decision Making*, 16(6), 1485–1505. <https://doi.org/10.1017/S1930297500008512>
- Zehr, J., & Schwarz, F. (2022). PennController for Internet Based Experiments (IBEX). *Open Science Framework*. <https://doi.org/10.17605/OSF.IO/MD832>
- Özsoy, O., Çiçek, B., Özal, Z., Gagarina, N., & Sekerina, I. A. (2023). Turkish-German heritage speakers' predictive use of case: webcam-based vs. in-lab eye-tracking. *Frontiers in Psychology*, 14, 1155585. <https://doi.org/10.3389/fpsyg.2023.1155585>